

PIPELINE PHASE BREAKOUT: What Each Deliverable Does and How It Fits

Lorespinner Pipeline Upgrade — May 2026

This document explains what each of the 9 deliverables accomplishes in plain language and shows exactly where it connects to the existing Logic Companion pipeline.

THE CURRENT PIPELINE (What Daniel Already Built)

The existing Logic Companion v1.1 runs in this order:

```
FORMAT DETECTION → PHASE 1 (IP Audit) → PHASE 2 (Story Session Map) →  
PHASE 3 (Entry Point) → PHASE 4 (Beat Architecture) → PHASE 5 (Choice Design) →  
PHASE 6 (Consequence Mapping) → PHASE 7 (Session Close) → PHASE 8 (Editorial  
Verification)
```

Plus the Chaos System Prompt (runtime narrator) which is hand-crafted per IP.

Everything below plugs into or upgrades this existing chain. Nothing is rebuilt from scratch.

THE UPGRADED PIPELINE (What It Looks Like After)

```
IP TRIMMING AGENT (new) → FORMAT DETECTION → PHASE 1 (IP Audit) →  
VOICE LOCK (new) → PHASE 2 (upgraded) → PHASE 3 (unchanged) →  
PHASE 4 (upgraded) → PHASE 5 (upgraded) → PHASE 6 (upgraded) →  
PHASE 7 (unchanged) → PHASE 8 (upgraded) →  
ASSEMBLY JOB (new) → RUNTIME NARRATOR TEMPLATE (replaces hand-crafted Chaos)
```

Post-experience: SOCIAL ECHO LAYER (new product feature)

DELIVERABLE 7: IP TRIMMING AGENT

Where it goes: Brand new job at the very front of the pipeline, before Format Detection.

What it does: Takes a full source IP (novel, screenplay, whatever) and surgically reduces it to its essential muscle. Strips excessive description, redundant scenes, backstory dumps, and exposition blocks while keeping every line of dialogue, every character-revealing action, every world rule, and every object or NPC that gets referenced later.

What it produces:

- A Story Spine (protagonist, dramatic question, turning points, climax)
- World Rules extraction (what can and cannot exist — feeds StoryGuard later)
- A content triage log (what was kept, what was cut, why)
- Conversion notes (how cut material should become interactive content — "this backstory dump becomes NPC dialogue," "this description becomes an explorable environment")
- The trimmed source text (25-45% smaller than the original)

How it works with the current logic: Every phase that currently receives the full source text now receives the trimmed version instead. This reduces token costs at every downstream phase. The one exception: the Voice Lock Phase gets the FULL original, because voice extraction needs the author's complete range.

What problem it solves: A 400-page novel is 100-120k tokens. After trimming, 60-70k. Over 12+ pipeline jobs, that savings compounds to 30-50% reduction in total processing cost. For big IPs like Lord of the Rings, this is the difference between feasible and impossible.

DELIVERABLE 1: VOICE LOCK PHASE

Where it goes: New phase that slots between Phase 1 (IP Audit) and Phase 2 (Story Session Map). Runs once per IP.

What it does: Forensically extracts the author's specific writing DNA so that every word the narrator speaks sounds like the author wrote it, not like AI imitating a writer.

What it produces:

- Author Voice DNA: 8-12 signature writing techniques with quoted evidence, sentence rhythm patterns, vocabulary fingerprint, dialogue fingerprints per character, emotional range map, paragraph architecture
- Master Rule 1 (Hard Ban List): Universal bans that apply to ALL stories (em dashes, AI vocabulary like "delve" and "tapestry," sentence molds like "It's not X, it's Y"), plus IP-specific bans generated from what THIS author never does
- 14-Point Audit Protocol: A quality gate that runs during Phase 8 editorial verification — not at runtime — to catch AI tells before they reach the player

How it works with the current logic: The current Chaos prompt already has voice markers and hard bans (the Alice prompt has them). Voice Lock STANDARDIZES this. Instead of hand-crafting voice rules per IP, the pipeline extracts them automatically from the source. The output feeds directly into the runtime narrator template (Sections 1, 3, and 4) and into Phase 8 (the voice audit questions).

What problem it solves: None of the three original stories sound like their actual author. DriftHeart sounds completely like AI writing. Words from Anima Machina are bleeding into Nocturne. Voice Lock creates an immune system that protects every IP's unique voice with the same rigor that StoryGuard protects canon.

DELIVERABLE 2: UPDATED PHASE 2 — STORY SESSION MAP

Where it goes: Replaces the existing Phase 2 prompt. Same position in the pipeline.

What stays: Everything the current Phase 2 already does — session allocation, arc progression, branch opportunities, cross-session payoff plan, branch dimension definitions. All five existing tasks are untouched.

What's added (4 new tasks appended):

Persistent State Schema — A master inventory of everything trackable in the world: every object the player can find or lose, every NPC with their disposition levels and behavioral triggers, every location with sensory details and environmental reactivity, emotional ledger categories (acts of mercy, betrayals, sacrifices, etc.), and action history categories (violent, deceptive, protective actions). This gives the runtime a concrete list to read and update every turn instead of improvising.

World Reactivity Rules — 5-8 trigger/response pairs defining how THIS specific world reacts to player behavior patterns. Not individual NPC reactions — world-level atmospheric shifts. "If Nora avoids emotional vulnerability, the technology around her glitches harder." "If Alice insists on logic, Wonderland escalates nonsense." Each rule is grounded in the IP's own world logic.

StoryGuard Canon Extraction — The four protection layers codified as concrete rules: what can/cannot exist in this world (Layer 1), narrative destinations that must be protected (Layer 2), per-character behavioral boundaries (Layer 3), and a scene rules template that Phase 5 fills per episode (Layer 4).

Story-Native Alignment Labels — The player-facing names for the chaotic/lawful/neutral alignment spectrum, using language that belongs in the story's world. "Disruptor / Architect / Ghost" for Anima Machina instead of RPG terminology.

What problem it solves: The "world forgot my knife" problem. The "NPC didn't react when I betrayed them" problem. The "creature that doesn't belong showed up" problem. Phase 2 now defines every trackable element and every boundary BEFORE any episode is designed.

DELIVERABLE 3: UPDATED PHASE 4 — BEAT ARCHITECTURE

Where it goes: Replaces the existing Phase 4 prompt. Same position in the pipeline.

What changes: The session structure is updated for 20-30 minute episodes with the recalibrated choice density model:

- 4 branching choices (load-bearing, forks the story, 115-125 word outcomes)
- 4-6 emotional choices (textural, colors narration, 80-100 word outcomes)
- 6-10 posture shifts (NEW — lightweight interactions within the narration flow, no stopping)

The five-beat arc is preserved (Setup, Escalation, Breath, Twist, Resolution) but the timing and interaction placement are updated. A new Episode Content Budget task calculates the word count per episode (target: 3,500-4,200 words) to prevent scope creep. A Posture Shift Placement task designs each in-narration interaction.

How it works with the current logic: The existing beat map format is preserved. Daniel's pipeline already knows how to process a beat map. The new version just has more interaction types mapped into it and a tighter content budget.

What problem it solves: The original 16-18 choices per 15-minute session was too aggressive — a choice every 50 seconds with no room to breathe. The new model gives 8-10 full choices plus 6-10 posture shifts that keep the player active without stopping the story.

DELIVERABLE 4: UPDATED PHASE 5 — CHOICE DESIGN

Where it goes: Replaces the existing Phase 5 prompt. Same position in the pipeline.

What stays: The core choice-writing rules — three options per choice, second-person present tense, one sentence per option, immediate downstream effects. All preserved.

What's added:

Chaotic/Lawful/Neutral Alignment System — Every branching and emotional choice maps its three options to chaotic (impulsive, boundary-breaking), lawful (measured, rule-following), and neutral (pragmatic, observational). NEVER labeled. NEVER in fixed order. Randomized per choice. Accumulates invisibly in player state over time.

StoryGuard Manifest per Choice Node — For every branching choice, a pre-planned manifest that tells the runtime exactly what to do when a player speaks their own choice instead of picking one of the three. Canon boundaries active at this moment, how each NPC would react to unexpected input, available objects from persistent state, and the nearest safe fold-back path.

Persistent State Changes per Choice — Every option specifies exactly what changes: inventory, NPC dispositions, environmental flags, emotional ledger entries, alignment shift.

World Noticed Signals — For every significant state change, a pre-authored in-world acknowledgment. Not "the world will remember this." Something physical: "Something behind Riven's eyes clicks shut. Not anger. Something quieter." Written in the author's voice.

Social Echo Defining Lines — Provocative one-liners per branching choice path for the post-episode share card.

What problem it solves: Freeform responses were improvised at runtime with no guardrails. Now every choice node has a pre-planned manifest so the runtime knows the boundaries before the player speaks. The alignment system creates subliminal replay value — two players have genuinely different experiences based on instinct alone.

DELIVERABLE 5: UPDATED PHASE 6 — CONSEQUENCE MAPPING

Where it goes: Replaces the existing Phase 6 prompt. Same position in the pipeline.

What stays: The existing four-column consequence table (Immediate Effect, Current Session Echo, Next Session Payoff, Later Session Legacy). Preserved exactly.

What's added (3 new columns per path):

World State Deltas — Exact inventory changes, NPC disposition shifts, environmental flag updates, alignment accumulation, and emotional ledger entries per choice path. This is the data layer that makes persistent state actually work.

Reactivity Triggers — Conditional world reactions that fire based on accumulated player state. "If player alignment is chaotic by this point, Riven reacts with wariness instead of warmth." "If player holds the knife from Episode 1, this scene changes."

Cross-Episode State Propagation Rules — What resets between episodes (temporary emotional states), what persists (inventory, relationships, alignment), and what escalates with accumulation (repeated behavioral patterns compound).

Freeform Consequence Guidelines — Per branching choice, guidelines for how the runtime maps freeform input to the nearest alignment and handles StoryGuard violations gracefully.

What problem it solves: Choices felt like they disappeared after one scene. The world didn't accumulate. Now every choice has explicit state changes that carry forward, and escalation rules that make repeated behavior patterns compound into visible world shifts.

DELIVERABLE 6: UPDATED PHASE 8 — EDITORIAL VERIFICATION

Where it goes: Replaces the existing Phase 8 prompt. Same position in the pipeline.

What stays: The original 10 design audit questions. Every one preserved exactly.

What's added (13 new questions in two new sections):

Voice Audit (7 questions): Hard ban scan (zero tolerance), trailing simile scan, sentence mold scan, voice authenticity test (could a reader identify the author?), full 14-point audit protocol execution, diction fingerprint check, and Social Echo defining lines verification.

StoryGuard and State Compliance (6 questions): Canon integrity (nothing outside the world), character truth (NPCs stay in character), world reactivity (the rules are visibly active), persistent state consistency (Phase 5 and Phase 6 data matches), World Noticed signals quality, alignment balance (three genuine options, not three versions of the same option).

Scoring: 23/23 = green light. 20-22 = amber (fix and re-run failed questions). Below 20 = red (return to flagged phases).

What problem it solves: The current 10-question gate catches design problems but misses voice drift and state inconsistencies. The upgraded gate catches AI tells, canon violations, and data mismatches before they reach the player.

DELIVERABLE 8: GENERALIZED RUNTIME NARRATOR TEMPLATE

Where it goes: Replaces hand-crafted Chaos prompts. An assembly job at the end of the pipeline takes all phase outputs and populates the template automatically.

What it is: A 17-section master template with a 65,000-character budget. Some sections are hardcoded (same for every IP — narration rules, freedom contract, mission statement). Some are populated by the pipeline (voice profile, world reference, character reference, session packet). Some are updated at runtime (persistent state, symbolic memory, arc position).

The 17 sections: Narrator Identity, World Reference, Character Reference, Voice Profile + Hard Bans, StoryGuard Layers, World Reactivity, Persistent World State (tiered loading), Symbolic Memory, Narrative Gravity, Current Arc Position, Session Packet (all choices + consequences + posture shifts), Full Episode Script, Cold Open, Narration Rules, Freedom Contract, Session-Complete Signal, Mission Statement.

Tiered State Loading: To prevent the prompt from ballooning as the player progresses — Tier 1 (always: current scene objects/NPCs/alignment), Tier 2 (scene-relevant: emotional ledger, connected locations), Tier 3 (episode transitions only: full history). Keeps runtime cost increase to 10-15%.

How it works with the current logic: Daniel builds one assembly job that reads all phase outputs and fills the template slots. That job replaces the current workflow of hand-crafting a Chaos prompt for each story. The Runtime Branch Resolution Policy from the existing Logic Companion is built into Section 15 (Freedom Contract).

What problem it solves: Hand-crafting a 60k-character Chaos prompt for every IP is not scalable. The template makes it automatic. Process any IP through the pipeline, the template assembles itself.

DELIVERABLE 9: SOCIAL ECHO LAYER SPECIFICATION

Where it goes: Product feature that ships after the core pipeline is working. This is a spec for the engineering and product team, not a pipeline prompt.

What it is: A post-episode social sharing system with four components:

Share Card — After each episode, the player gets a shareable card showing their alignment tendency as a visual spectrum, one provocative line about their defining choice (in the author's voice), and a comparison stat ("47% of players made the same choice").

Alignment Profile — Across episodes, the player builds a cumulative identity using story-native language ("Disruptor / Architect / Ghost"), visible only post-episode, evolving over time.

Comparison Hooks — Aggregate player data surfaced post-episode: rarity stats, correlation data, friend comparisons. Never shown during gameplay. Never influences the experience.

Replay Invitation — Subtle hooks: "There were 8 moments in this episode you never saw." "Your story diverged from [Friend]'s at the gate." Creates replay desire through curiosity, not instruction.

What problem it solves: Zero social = zero virality. The interactive story market lives on word-of-mouth and social sharing. This turns every completed episode into a marketing moment.

HOW IT ALL CONNECTS

The pipeline is a data chain. Each phase produces output that the next phase consumes:

- IP Trimming Agent produces the trimmed source + world rules + conversion notes
- Voice Lock produces the voice DNA + ban list + audit protocol

- Phase 2 produces the session map + persistent state schema + world reactivity + StoryGuard + alignment labels
- Phase 4 produces the beat map + content budget + posture shift placements
- Phase 5 produces all choices with alignment + state changes + StoryGuard manifests + World Noticed signals + defining lines
- Phase 6 produces consequence maps + state deltas + reactivity triggers + propagation rules
- Phase 8 runs the 23-question quality gate
- The Assembly Job takes ALL of the above and populates the Runtime Template
- The Runtime Template is the complete system prompt the AI narrator receives at game time
- Social Echo reads the runtime state data post-episode and generates share cards

Break any link in the chain and the downstream phases lose data. That's why Phase 2 through Phase 8 must be deployed together.

END OF PHASE BREAKOUT